

QTW – Module 13
Preession
Jessica McPhaul

3 Key Takeaways:

1. RNNs process sequences step-by-step and retain context by feeding each output into the next input.
2. Padding, embedding, and fixed-length preprocessing are essential to make language data usable by RNNs.
3. LSTM and GRU improve memory retention over long sequences and are often better than basic RNNs for NLP.

3 Questions:

1. How do we decide between LSTM and GRU for practical NLP tasks in terms of training time and performance?
2. When is it more effective to use sequence output versus the final timestep in RNN-based classification?
3. What are the trade-offs between using pretrained embeddings versus learning them within the model?