

```
# prompt: mount drive

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

import psutil
print(f"Available Memory: {psutil.virtual_memory().available / 1e9:.2f} GB")

Available Memory: 86.01 GB

import torch
import cupy as cp

# Check PyTorch CUDA availability
print(f"PyTorch CUDA available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"PyTorch Device: {torch.cuda.get_device_name(0)}")
    print(f"CUDA Version (PyTorch): {torch.version.cuda}")

# Check CuPy CUDA availability
print(f"CuPy CUDA available: {cp.cuda.is_available()}")
if cp.cuda.is_available():
    print(f"CUDA Version (CuPy): {cp.cuda.runtime.runtimeGetVersion() / 1000}")

import torch

if torch.cuda.is_available():
    print("CUDA is available!")
    print("Device:", torch.cuda.get_device_name(0))
else:
    print("CUDA is NOT available.")

import cudf

print("cudf is successfully installed!")
df = cudf.DataFrame({'a': [1, 2, 3], 'b': [4, 5, 6]})
print(df)

PyTorch CUDA available: True
PyTorch Device: NVIDIA A100-SXM4-40GB
CUDA Version (PyTorch): 12.4
CuPy CUDA available: True
CUDA Version (CuPy): 12.06
CUDA is available!
Device: NVIDIA A100-SXM4-40GB
cudf is successfully installed!
   a  b
0  1  4
1  2  5
2  3  6

# 2. EDA

import cudf

# Load the data into cudf DataFrames
diabetic_data = cudf.read_csv("/content/drive/MyDrive/diabetic_data.csv")
ids_mapping = cudf.read_csv("/content/drive/MyDrive/IDs_mapping.csv")

# Ensure all string columns are treated as string type
diabetic_data = diabetic_data.astype(str)

# Replace '?' with None before converting to cudf's NA
diabetic_data = diabetic_data.replace({'?': None}).fillna(cudf.NA)

# or (if needed) fix - convert only object columns
# for col in diabetic_data.select_dtypes(include='object'):
#     diabetic_data[col] = diabetic_data[col].replace({'?': None}).fillna(cudf.NA)

# Display dataset info
print("\n Diabetic Data Info:")
print(diabetic_data.info())

print("\n First few rows of diabetic_data:")
print(diabetic_data.head())

print("\n IDs Mapping Data Info:")
print(ids_mapping.info())

print("\n First few rows of IDs_mapping:")
print(ids_mapping.head())

# Check missing values
print("\n Missing values in dataset:")
missing_counts = diabetic_data.isnull().sum()
```

```
print(missing_counts[missing_counts > 0])

# 2.b
# # Convert columns back to proper types:
for col in diabetic_data.columns:
    if diabetic_data[col].str.isnumeric().all():
        diabetic_data[col] = diabetic_data[col].astype("int64")

Diabetic Data Info:
<class 'cudf.core.dataframe.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 50 columns):
#   Column                                Non-Null Count  Dtype
---  ---
0   encounter_id                          101766 non-null object
1   patient_nbr                           101766 non-null object
2   race                                  99493 non-null object
3   gender                               101766 non-null object
4   age                                   101766 non-null object
5   weight                               3197 non-null object
6   admission_type_id                    101766 non-null object
7   discharge_disposition_id             101766 non-null object
8   admission_source_id                 101766 non-null object
9   time_in_hospital                    101766 non-null object
10  payer_code                           61510 non-null object
11  medical_specialty                    51817 non-null object
12  num_lab_procedures                  101766 non-null object
13  num_procedures                      101766 non-null object
14  num_medications                     101766 non-null object
15  number_outpatient                   101766 non-null object
16  number_emergency                    101766 non-null object
17  number_inpatient                    101766 non-null object
18  diag_1                              101745 non-null object
19  diag_2                              101408 non-null object
20  diag_3                              100343 non-null object
21  number_diagnoses                    101766 non-null object
22  max_glu_serum                       101766 non-null object
23  A1cresult                           101766 non-null object
24  metformin                           101766 non-null object
25  repaglinide                         101766 non-null object
26  nateglinide                         101766 non-null object
27  chlorpropamide                     101766 non-null object
28  glimepiride                         101766 non-null object
29  acetohexamide                     101766 non-null object
30  glipizide                           101766 non-null object
31  glyburide                           101766 non-null object
32  tolbutamide                         101766 non-null object
33  pioglitazone                        101766 non-null object
34  rosiglitazone                       101766 non-null object
35  acarbose                            101766 non-null object
36  miglitol                            101766 non-null object
37  troglitazone                       101766 non-null object
38  tolazamide                          101766 non-null object
39  examide                             101766 non-null object
40  citoglipton                         101766 non-null object
41  insulin                             101766 non-null object
42  glyburide-metformin                 101766 non-null object
43  glipizide-metformin                 101766 non-null object
44  glimepiride-pioglitazone            101766 non-null object
45  metformin-rosiglitazone             101766 non-null object
46  metformin-pioglitazone              101766 non-null object
47  change                              101766 non-null object
48  diabetesMed                         101766 non-null object
49  readmitted                         101766 non-null object
dtypes: object(50)

# Step 2.2 Fix Data Types and Handle Missing Values

import cudf

# Convert Numeric Columns First
numeric_cols = [
    "encounter_id", "patient_nbr", "admission_type_id", "discharge_disposition_id",
    "admission_source_id", "time_in_hospital", "num_lab_procedures", "num_procedures",
    "num_medications", "number_outpatient", "number_emergency", "number_inpatient",
    "number_diagnoses"
]

for col in numeric_cols:
    diabetic_data[col] = diabetic_data[col].astype("int64")

# Convert Categorical Columns to String and Replace Missing Values
categorical_cols = [
    "race", "gender", "age", "payer_code", "medical_specialty",
    "diag_1", "diag_2", "diag_3", "max_glu_serum", "A1cresult",
    "metformin", "repaglinide", "nateglinide", "chlorpropamide",
    "glimepiride", "acetohexamide", "glipizide", "glyburide",
]
```

```
"tolbutamide", "pioglitazone", "rosiglitazone", "acarbose",
"miglitol", "troglitazone", "tolazamide", "examide",
"citoglipton", "insulin", "glyburide-metformin",
"glipizide-metformin", "glimepiride-pioglitazone",
"metformin-rosiglitazone", "metformin-pioglitazone",
"change", "diabetesMed", "readmitted"
]

for col in categorical_cols:
    diabetic_data[col] = diabetic_data[col].astype("str").replace({'?': cudf.NA})

# Verify Fix
print(" Data Types Fixed and Missing Values Handled!")
print(diabetic_data.dtypes)
```

🔗 Data Types Fixed and Missing Values Handled!

encounter_id	int64
patient_nbr	int64
race	object
gender	object
age	object
weight	object
admission_type_id	int64
discharge_disposition_id	int64
admission_source_id	int64
time_in_hospital	int64
payer_code	object
medical_specialty	object
num_lab_procedures	int64
num_procedures	int64
num_medications	int64
number_outpatient	int64
number_emergency	int64
number_inpatient	int64
diag_1	object
diag_2	object
diag_3	object
number_diagnoses	int64
max_glu_serum	object
AlCresult	object
metformin	object
repaglinide	object
nateglinide	object
chlorpropamide	object
glimepiride	object
acetohexamide	object
glipizide	object
glyburide	object
tolbutamide	object
pioglitazone	object
rosiglitazone	object
acarbose	object
miglitol	object
troglitazone	object
tolazamide	object
examide	object
citoglipton	object
insulin	object
glyburide-metformin	object
glipizide-metformin	object
glimepiride-pioglitazone	object
metformin-rosiglitazone	object
metformin-pioglitazone	object
change	object
diabetesMed	object
readmitted	object
dtype:	object

```
# Merge ids
import cudf

# Check for non-numeric values
invalid_values = ids_mapping[ids_mapping["admission_type_id"].str.isnumeric()]
print(" Non-Numeric Values in 'admission_type_id':", invalid_values)

# Convert numeric values to integers
ids_mapping = ids_mapping[ids_mapping["admission_type_id"].str.isnumeric()]
ids_mapping["admission_type_id"] = ids_mapping["admission_type_id"].astype("int64")

print("\n Cleaned 'ids_mapping' Data:")
print(ids_mapping.head())
```

🔗 Non-Numeric Values in 'admission_type_id':

	admission_type_id	description
9	discharge_disposition_id	description
41	admission_source_id	description

Cleaned 'ids_mapping' Data:

	admission_type_id	description
0	1	Emergency
1	2	Urgent
2	3	Elective
3	4	Newborn
4	5	Not Available

```
# 3.2

# Merge 'diabetic_data' with 'ids_mapping' on 'admission_type_id'
diabetic_data = diabetic_data.merge(ids_mapping, how="left", on="admission_type_id")

# Drop unnecessary columns
columns_to_drop = [
    "weight", "max_glu_serum", "AlCresult", "medical_specialty", "payer_code",
    "encounter_id", "patient_nbr", "description" # 'description' is from ids_mapping
]
diabetic_data = diabetic_data.drop(columns=columns_to_drop)

# Fill Missing Values in Key Categorical Columns
for col in ["race", "diag_1", "diag_2", "diag_3"]:
    diabetic_data[col] = diabetic_data[col].fillna("Unknown")

# Convert 'readmitted' to numerical categories
diabetic_data["readmitted"] = diabetic_data["readmitted"].map({"NO": 0, ">30": 1, "<30": 2})

# Verify Merge & Cleaning
print(" Merge Completed and Data Cleaned!")
print(diabetic_data.dtypes)
print("\n First Few Rows of Cleaned Data:")
print(diabetic_data.head())
```

🔗 Merge Completed and Data Cleaned!

race	object
gender	object
age	object
admission_type_id	int64
discharge_disposition_id	int64
admission_source_id	int64
time_in_hospital	int64
num_lab_procedures	int64
num_procedures	int64
num_medications	int64
number_outpatient	int64
number_emergency	int64
number_inpatient	int64
diag_1	object
diag_2	object
diag_3	object
number_diagnoses	int64
metformin	object
repaglinide	object
nateglinide	object
chlorpropamide	object
glimepiride	object
acetohexamide	object
glipizide	object
glyburide	object
tolbutamide	object
pioglitazone	object
rosiglitazone	object
acarbose	object
miglitol	object
troglitazone	object
tolazamide	object
examide	object
citoglipton	object
insulin	object
glyburide-metformin	object
glipizide-metformin	object
glimepiride-pioglitazone	object
metformin-rosiglitazone	object
metformin-pioglitazone	object
change	object
diabetesMed	object
readmitted	int64
dtype:	object

First Few Rows of Cleaned Data:

	race	gender	age	admission_type_id	\
0	AfricanAmerican	Female	[70-80)	3	
1	AfricanAmerican	Female	[70-80)	3	
2	AfricanAmerican	Female	[70-80)	3	
3	AfricanAmerican	Male	[50-60)	2	
4	AfricanAmerican	Male	[50-60)	2	

```

    discharge_disposition_id  admission_source_id  time_in_hospital  \
0                             3                    2                6
1                             3                    2                6
2                             3                    2                6

import pandas as pd
from sklearn.preprocessing import StandardScaler

# Load the datasets
diabetic_data = pd.read_csv("../content/drive/MyDrive/MSL_Case Study 2/diabetic_data.csv")
ids_mapping = pd.read_csv("../content/drive/MyDrive/MSL_Case Study 2/IDs_mapping.csv")

# 3.2 Data Cleaning and Merging
# Convert 'admission_type_id' to numeric, handling non-numeric values
diabetic_data['admission_type_id'] = pd.to_numeric(diabetic_data['admission_type_id'], errors='coerce')
ids_mapping['admission_type_id'] = pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce')

# Convert to Int64 after ensuring both are numeric
diabetic_data['admission_type_id'] = diabetic_data['admission_type_id'].astype('Int64')
ids_mapping['admission_type_id'] = ids_mapping['admission_type_id'].astype('Int64')

# Merge diabetic_data with ids_mapping (now with consistent data types)
diabetic_data = diabetic_data.merge(ids_mapping, how='left', on='admission_type_id')

# Fill missing values in key categorical columns
for col in ['race', 'diag_1', 'diag_2', 'diag_3']:
    diabetic_data[col] = diabetic_data[col].fillna('Unknown')

# Convert 'readmitted' to numerical categories
diabetic_data['readmitted'] = diabetic_data['readmitted'].map({'NO': 0, 'YES': 1, 'LARGE': 2})

# Convert 'max_glu_serum' and 'A1cresult' to numerical representations
diabetic_data['max_glu_serum'] = diabetic_data['max_glu_serum'].replace({
    'None': 0,
    'None': 1,
    '>200': 2,
    '>300': 3
})

diabetic_data['A1cresult'] = diabetic_data['A1cresult'].replace({
    'None': 0,
    'None': 1,
    '>7': 2,
    '>8': 3
})

# 4. Feature Engineering (Scaling Numeric Features)

# Define Numeric Columns
numeric_cols = [
    "time_in_hospital", "num_lab_procedures", "num_procedures",
    "num_medications", "number_outpatient", "number_emergency",
    "number_inpatient", "number_diagnoses"
]

# Initialize StandardScaler
scaler = StandardScaler()

# Fit and transform the selected numeric columns
diabetic_data[numeric_cols] = scaler.fit_transform(diabetic_data[numeric_cols])

# Verify Merge, Cleaning, and Scaling
print("Merge Completed and Data Cleaned!")
print(diabetic_data.dtypes)
print("\nFirst Few Rows of Cleaned Data:")
print(diabetic_data.head())
```

🔗 Merge Completed and Data Cleaned!

encounter_id	int64
patient_nbr	int64
race	object
gender	object
age	object
weight	object
admission_type_id	Int64
discharge_disposition_id	int64
admission_source_id	int64
time_in_hospital	float64
payer_code	object
medical_specialty	object
num_lab_procedures	float64
num_procedures	float64
num_medications	float64
number_outpatient	float64
number_emergency	float64
number_inpatient	float64
diag_1	object
diag_2	object
diag_3	object
number_diagnoses	float64

max_glu_serum	float64
A1cresult	float64
metformin	object
repaglinide	object
nateglinide	object
chlorpropamide	object
glimepiride	object
acetohexamide	object
glipizide	object
glyburide	object
tolbutamide	object
pioglitazone	object
rosiglitazone	object
acarbose	object
miglitol	object
troglitazone	object
tolazamide	object
examide	object
citoglipton	object
insulin	object
glyburide-metformin	object
glipizide-metformin	object
glimepiride-pioglitazone	object
metformin-rosiglitazone	object
metformin-pioglitazone	object
change	object
diabetesMed	object
readmitted	int64
description	object
dtype: object	

First Few Rows of Cleaned Data:

encounter_id	patient_nbr	race	gender	age	weight	
1	1					

```
# Save the cleaned data to 'data_cleaned.csv'
diabetic_data.to_csv("data_cleaned.csv", index=False)
```

```
# Fix data types and handle missing values
```

```
from google.colab import drive
import psutil
import torch
import cupy as cp
import pandas as pd
from sklearn.preprocessing import StandardScaler

drive.mount('/content/drive')

print(f"Available Memory: (psutil.virtual_memory().available / 1e9:.2f) GB")

# Check PyTorch CUDA availability
print(f"PyTorch CUDA available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"PyTorch Device: {torch.cuda.get_device_name(0)}")
    print(f"CUDA Version (PyTorch): {torch.version.cuda}")

# Check CuPy CUDA availability
print(f"CuPy CUDA available: {cp.cuda.is_available()}")
if cp.cuda.is_available():
    print(f"CUDA Version (CuPy): {cp.cuda.runtime.runtimeGetVersion() / 1000}")

if torch.cuda.is_available():
    print("CUDA is available")
    print("Device:", torch.cuda.get_device_name(0))
else:
    print("CUDA is NOT available.")

print("cuDF is successfully installed!") #This line seems unnecessary, remove it if you don't need to confirm installation
```

```
# Load the datasets using pandas
diabetic_data = pd.read_csv("../content/drive/MyDrive/MSL_Case Study 2/diabetic_data.csv")
ids_mapping = pd.read_csv("../content/drive/MyDrive/MSL_Case Study 2/IDs_mapping.csv")

# 3.2 Data Cleaning and Merging
# Convert 'admission_type_id' to numeric, handling non-numeric values
diabetic_data['admission_type_id'] = pd.to_numeric(diabetic_data['admission_type_id'], errors='coerce')
ids_mapping['admission_type_id'] = pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce')

# Convert to Int64 after ensuring both are numeric
diabetic_data['admission_type_id'] = diabetic_data['admission_type_id'].astype('Int64')
ids_mapping['admission_type_id'] = ids_mapping['admission_type_id'].astype('Int64')

# Merge diabetic_data with ids_mapping (now with consistent data types)
diabetic_data = diabetic_data.merge(ids_mapping, how='left', on='admission_type_id')

# Fill missing values in key categorical columns
for col in ['race', 'diag_1', 'diag_2', 'diag_3']:
    diabetic_data[col] = diabetic_data[col].fillna('Unknown')
```

```
# Convert 'readmitted' to numerical categories
diabetic_data['readmitted'] = diabetic_data['readmitted'].map({"NO": 0, ">30": 1, "<=30": 2})

# Convert 'max_glu_serum' and 'A1Cresult' to numerical representations
diabetic_data['max_glu_serum'] = diabetic_data['max_glu_serum'].replace({
    'None': 0,
    'Norm': 1,
    '>200': 2,
    '>300': 3
})

diabetic_data['A1Cresult'] = diabetic_data['A1Cresult'].replace({
    'None': 0,
    'Norm': 1,
    '>7': 2,
    '>8': 3
})

# 4. Feature Engineering (Scaling Numeric Features)

# Define Numeric Columns
numeric_cols = [
    "time_in_hospital", "num_lab_procedures", "num_procedures",
    "num_medications", "number_outpatient", "number_emergency",
    "number_inpatient", "number_diagnoses"
]

# Initialize StandardScaler
scaler = StandardScaler()

# Fit and transform the selected numeric columns
diabetic_data[numeric_cols] = scaler.fit_transform(diabetic_data[numeric_cols])

# Verify Merge, Cleaning, and Scaling
print("Merge Completed and Data Cleaned!")
print(diabetic_data.dtypes)
print("\nFirst Few Rows of Cleaned Data:")
print(diabetic_data.head())

# Save the cleaned data to 'data_cleaned.csv'
diabetic_data.to_csv("data_cleaned.csv", index=False)

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
Available Memory: 85.14 GB
PyTorch CUDA available: True
PyTorch Device: NVIDIA A100-SXM4-40GB
CUDA Version (PyTorch): 12.4
CuPy CUDA available: True
CUDA Version (CuPy): 12.06
CUDA is available!
Device: NVIDIA A100-SXM4-40GB
cuDF is successfully installed!
<ipython-input-18-e075a01008e8>:58: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future ver
diabetic_data['max_glu_serum'] = diabetic_data['max_glu_serum'].replace({
<ipython-input-18-e075a01008e8>:65: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future ver
diabetic_data['A1Cresult'] = diabetic_data['A1Cresult'].replace({
Merge Completed and Data Cleaned!
encounter_id      int64
patient_nbr       int64
race              object
gender            object
age              object
weight            object
admission_type_id  int64
discharge_disposition_id  int64
admission_source_id  int64
time_in_hospital  float64
payer_code        object
medical_specialty object
num_lab_procedures float64
num_procedures    float64
num_medications   float64
number_outpatient float64
number_emergency  float64
number_inpatient  float64
diag_1            object
diag_2            object
diag_3            object
number_diagnoses  float64
max_glu_serum     float64
A1Cresult         float64
metformin         object
repaglinide       object
nateglinide       object
chlorpropamide    object
glinepiride       object
acetohexamide     object
glipizide         object
glyburide         object
```

```
tolbutamide      object
pioglitazone     object
rosiglitazone    object
acarbose         object
miglitol         object
troglitazone     object
tolazamide       object
examide          object
citoglipton      object
insulin          object

# Drop unnecessary columns
columns_to_drop = [
    "weight", "max_glu_serum", "A1Cresult", "medical_specialty", "payer_code",
    "encounter_id", "patient_nbr", "description" # 'description' is from ids_mapping
]
diabetic_data = diabetic_data.drop(columns=columns_to_drop, errors='ignore') # Use errors='ignore'

# Check for non-numeric values and handle them

# Check if 'admission_type_id' is numeric using pd.to_numeric
invalid_values = ids_mapping[pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce').isnull()]

if not invalid_values.empty:
    print("Non-Numeric Values in 'admission_type_id':\n", invalid_values)
    # Decide how to handle invalid values: remove them, convert to numeric, or fill with a specific value
    # Option 1: Remove rows with non-numeric values
    ids_mapping = ids_mapping[ids_mapping['admission_type_id'].str.isnumeric()] # str is not needed here
    ids_mapping = ids_mapping[pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce').notnull()]

    # Option 2: Convert non-numeric values to a default numeric value
    ids_mapping.loc[ids_mapping['admission_type_id'].str.isnumeric(), 'admission_type_id'] = 0 # Example: replace with 0 # str is not needed here
    ids_mapping.loc[pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce').isnull(), 'admission_type_id'] = 0 # Example: replace with 0

# Convert numeric values to integers

# Print cleaned data
print("\nCleaned 'ids_mapping' Data:")
print(ids_mapping.head())

Non-Numeric Values in 'admission_type_id':
      admission_type_id  description
8                    <NA>      NaN
9                    <NA>      NaN
40                   <NA>      NaN
41                   <NA>      NaN

Cleaned 'ids_mapping' Data:
      admission_type_id  description
0                      1      Emergency
1                      2      Urgent
2                      3      Elective
3                      4      Newborn
4                      5      Not Available

# Check for non-numeric values and handle them
invalid_values = ids_mapping[pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce').isnull()]

if not invalid_values.empty:
    print("Non-Numeric Values in 'admission_type_id':\n", invalid_values)
    # Remove rows with non-numeric values
    ids_mapping = ids_mapping[pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce').notnull()]

# Convert 'admission_type_id' to numeric in both DataFrames
ids_mapping['admission_type_id'] = pd.to_numeric(ids_mapping['admission_type_id'], errors='coerce').astype('Int64')
diabetic_data['admission_type_id'] = pd.to_numeric(diabetic_data['admission_type_id'], errors='coerce').astype('Int64')

# Merge the DataFrames
diabetic_data = diabetic_data.merge(ids_mapping, how='left', on='admission_type_id')

# Drop unnecessary columns
columns_to_drop = [
    "weight", "max_glu_serum", "A1Cresult", "medical_specialty", "payer_code",
    "encounter_id", "patient_nbr", "description" # 'description' is from ids_mapping
]
diabetic_data = diabetic_data.drop(columns=columns_to_drop, errors='ignore') # Use errors='ignore'

# Fill Missing Values in Key Categorical Columns
for col in ["race", "diag_1", "diag_2", "diag_3"]:
    diabetic_data[col] = diabetic_data[col].fillna("Unknown")

# Convert 'readmitted' to numerical categories
diabetic_data['readmitted'] = diabetic_data['readmitted'].map({"NO": 0, ">30": 1, "<=30": 2})

# Verify Merge & Cleaning
print(" Merge Completed and Data Cleaned!")
```

```
print(diabetic_data.dtypes)
print("\n First Few Rows of Cleaned Data:")
print(diabetic_data.head())
```

```

Merge Completed and Data Cleaned!
race          object
gender        object
age           object
admission_type_id    int64
discharge_disposition_id  int64
admission_source_id    int64
time_in_hospital  float64
num_lab_procedures  float64
num_procedures    float64
num_medications    float64
number_outpatient  float64
number_emergency   float64
number_inpatient   float64
diag_1           object
diag_2           object
diag_3           object
number_diagnoses  float64
metformin        object
repaglinide       object
nateglinide       object
chlorpropamide    object
glimepiride       object
acetohexamide     object
glipizide         object
glyburide         object
tolbutamide       object
pioglitazone      object
rosiglitazone     object
acarbose          object
miglitol          object
troglitazone      object
tolazamide        object
examide           object
citoglipton       object
insulin           object
glyburide-metformin  object
glipizide-metformin  object
glimepiride-pioglitazone  object
metformin-rosiglitazone  object
metformin-pioglitazone  object
change           object
diabetesMed       object
readmitted        float64
dtype: object

First Few Rows of Cleaned Data:
   race gender age admission_type_id discharge_disposition_id \
0  Caucasian Female [0-10]          6                25
1  Caucasian Female [0-10]          6                25
2  Caucasian Female [0-10]          6                25
3  Caucasian Female [0-10]          6                25
4  Caucasian Female [0-10]          6                25

   admission_source_id time_in_hospital num_lab_procedures num_procedures \
0                    1          -1.137649          -0.106517          -0.785398
1                    1          -1.137649          -0.106517          -0.785398
2                    1          -1.137649          -0.106517          -0.785398
```

```

import pandas as pd

# df
categorical_cols = ["race", "gender", "age", "change", "diabetesMed", "insulin"]

# Use pandas get_dummies for one-hot encoding
diabetic_data = pd.get_dummies(diabetic_data, columns=categorical_cols, dummy_na=True)

print("Categorical Features One-Hot Encoded Successfully!")
print(diabetic_data.head())
```

```

Categorical Features One-Hot Encoded Successfully!
admission_type_id discharge_disposition_id admission_source_id \
0                6                25                1
1                6                25                1
2                6                25                1
3                6                25                1
4                6                25                1

time_in_hospital num_lab_procedures num_procedures num_medications \
0          -1.137649          -0.106517          -0.785398          -1.848268
1          -1.137649          -0.106517          -0.785398          -1.848268
```

```

2          -1.137649          -0.106517          -0.785398          -1.848268
3          -1.137649          -0.106517          -0.785398          -1.848268
4          -1.137649          -0.106517          -0.785398          -1.848268

number_outpatient number_emergency number_inpatient ... change_No \
0          -0.291461          -0.21262          -0.503276 ...      True
1          -0.291461          -0.21262          -0.503276 ...      True
2          -0.291461          -0.21262          -0.503276 ...      True
3          -0.291461          -0.21262          -0.503276 ...      True
4          -0.291461          -0.21262          -0.503276 ...      True
```

```

change_nan diabetesMed_No diabetesMed_Yes diabetesMed_nan insulin_Down \
0      False      True      False      False      False
1      False      True      False      False      False
2      False      True      False      False      False
3      False      True      False      False      False
4      False      True      False      False      False

insulin_No insulin_Steady insulin_Up insulin_nan
0      True      False      False      False
1      True      False      False      False
2      True      False      False      False
3      True      False      False      False
4      True      False      False      False
```

[5 rows x 70 columns]

```

# print(diabetic_data.head())

import pandas as pd

# Define categorical columns (This line might be redundant if already defined)
categorical_cols = ["race", "gender", "age", "change", "diabetesMed", "insulin"]

# Check if columns exist before applying get_dummies
if all(col in diabetic_data.columns for col in categorical_cols):
    # Use pandas get_dummies for one-hot encoding if columns are present
    diabetic_data = pd.get_dummies(diabetic_data, columns=categorical_cols, dummy_na=True)
    print("Categorical Features One-Hot Encoded Successfully!")
    print(diabetic_data.head())
else:
    print("Categorical columns have already been encoded or do not exist in the DataFrame.")
```

Categorical columns have already been encoded or do not exist in the DataFrame.

```

# Convert 'diag_1', 'diag_2', 'diag_3' to categorical codes
for col in ['diag_1', 'diag_2', 'diag_3']:
    diabetic_data[col] = diabetic_data[col].astype('category').cat.codes

# Convert all medication columns to binary (0/1)
medication_cols = [
    'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide', 'glimepiride',
    'acetohexamide', 'glipizide', 'glyburide', 'tolbutamide', 'pioglitazone',
    'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone', 'tolazamide',
    'examide', 'citoglipton', 'glyburide-metformin', 'glipizide-metformin',
    'glimepiride-pioglitazone', 'metformin-rosiglitazone', 'metformin-pioglitazone'
]

for col in medication_cols:
    # Convert only if the column is of string type
    if diabetic_data[col].dtype == 'object':
        diabetic_data[col] = (diabetic_data[col].astype(str) != "No").astype('int32')

# Drop the 'description' column if it exists
if 'description' in diabetic_data.columns:
    diabetic_data.drop(columns=['description'], inplace=True)

# Convert everything to float32
diabetic_data = diabetic_data.astype('float32')
print("All Features Converted to Numeric Format!")

print(diabetic_data['readmitted'].dtype)
print(diabetic_data['readmitted'].unique())
non_numeric_cols = diabetic_data.drop(columns=['readmitted']).select_dtypes(exclude=['number']).columns
print("Non-Numeric Columns in X:", non_numeric_cols)
```

All Features Converted to Numeric Format!
float32
[nan]
Non-Numeric Columns in X: Index([], dtype='object')

```

from sklearn.model_selection import train_test_split

# Define Features (X) and Target (y)
X = diabetic_data.drop(columns=['readmitted'])
# Convert to int32 and handle non-finite values with fillna
y = diabetic_data['readmitted'].fillna(-1).astype('int32') # Replace NaN with -1 before conversion

# Split Data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)
```

```
print(" Train/Test Split Completed! Shapes:")
print(f" - X_train: {X_train.shape}, y_train: {y_train.shape}")
print(f" - X_test: {X_test.shape}, y_test: {y_test.shape}")
```

Train/Test Split Completed! Shapes:

- X_train: (732715, 69), y_train: (732715,)
- X_test: (183179, 69), y_test: (183179,)

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.impute import SimpleImputer # Import SimpleImputer
```

```
# Load the cleaned data
diabetic_data = pd.read_csv("data_cleaned.csv")
```

```
# Define Features (X) and Target (y)
X = diabetic_data.drop(columns=['readmitted'])
y = diabetic_data['readmitted'].astype("int32")
```

```
# Handle potential non-numeric columns in X
non_numeric_cols = X.select_dtypes(exclude=['number']).columns
if not non_numeric_cols.empty:
    print("Warning: Non-numeric columns found in X:", non_numeric_cols)
    # Decide how to handle them (e.g., one-hot encoding, dropping)
    X = X.select_dtypes(include=['number'])
```

```
# Impute missing values using SimpleImputer
imputer = SimpleImputer(strategy='mean') # or 'median', 'most_frequent'
X = imputer.fit_transform(X) # Fit and transform to replace NAs
```

```
# Split Data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)
```

```
# Initialize and Train Model
log_reg = LogisticRegression(max_iter=1000, tol=1e-4)
log_reg.fit(X_train, y_train)
print("Logistic Regression Model Trained Successfully!")
```

```
# Predict on Test Data
y_pred = log_reg.predict(X_test)
```

```
# Compute Evaluation Metrics
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
```

```
# Display Results
print(f"Accuracy: {accuracy:.4f}")
print("\nConfusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)
```

```
# Check Class Imbalance
print("Class Distribution in Training Data:")
print(y_train.value_counts())
print("Class Distribution in Testing Data:")
print(y_test.value_counts())
```

Warning: Non-numeric columns found in X: Index(['race', 'gender', 'age', 'weight', 'payer_code', 'medical_specialty', 'diag_1', 'diag_2', 'diag_3', 'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide', 'glimepiride', 'acetohexamide', 'glipizide', 'glyburide', 'tolbutamide', 'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone', 'tolazamide', 'examide', 'cigtoglipton', 'insulin', 'glyburide-metformin', 'glipizide-metformin', 'glimepiride-pioglitazone', 'metformin-rosiglitazone', 'metformin-pioglitazone', 'change', 'diabetesMed', 'description'], dtype='object')

Logistic Regression Model Trained Successfully!
Accuracy: 0.5422

Confusion Matrix:

[[30102	2817	0]
[18320	3007	0]
[5974	840	0]]

Classification Report:		precision	recall	f1-score	support
	0	0.55	0.91	0.69	32919
	1	0.45	0.14	0.21	21327
	2	0.00	0.00	0.00	6814
	accuracy			0.54	61060
	macro avg	0.33	0.35	0.30	61060
	weighted avg	0.46	0.54	0.45	61060

Class Distribution in Training Data:

readmitted

0 131673

```
1 85308
2 27257
Name: count, dtype: int64
Class Distribution in Testing Data:
readmitted
0 32919
1 21327
2 6814
Name: count, dtype: int64
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

#visualize results and provide analysis

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# ... (your existing code) ...
```

```
# Display Results
print(f"Accuracy: {accuracy:.4f}")
print("\nConfusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)
```

```
# Visualize the Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
            xticklabels=["No Readmission", "Readmitted >30", "Readmitted <30"],
            yticklabels=["No Readmission", "Readmitted >30", "Readmitted <30"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
# Analyze Class Distribution
plt.figure(figsize=(6, 4))
sns.countplot(x=y_train) # or y_test
plt.title("Class Distribution")
plt.xlabel("Readmission Category")
plt.ylabel("Number of Patients")
plt.show()
```

```
# Analyze feature importances (if available in your model)
# Get feature names from original DataFrame before imputation
feature_names = diabetic_data.drop(columns=['readmitted']).columns
```

prompt: visualize results and provide analysis

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# ... (your existing code) ...
```

```
# Display Results
print(f"Accuracy: {accuracy:.4f}")
print("\nConfusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)
```

```
# visualize the Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
            xticklabels=["No Readmission", "Readmitted >30", "Readmitted <30"],
            yticklabels=["No Readmission", "Readmitted >30", "Readmitted <30"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
# Analyze Class Distribution
plt.figure(figsize=(6, 4))
sns.countplot(x=y_train) # or y_test
plt.title("Class Distribution")
plt.xlabel("Readmission Category")
plt.ylabel("Number of Patients")
plt.show()
```



Accuracy: 0.5422

Confusion Matrix:

		0	1	2
0	30102	2817	0	
1	18320	3007	0	
2	5974	840	0	

Classification Report:

	precision	recall	f1-score	support
0	0.55	0.91	0.69	32919
1	0.45	0.14	0.21	21327
2	0.00	0.00	0.00	6814
accuracy	0.33	0.35	0.54	61060
macro avg	0.33	0.35	0.30	61060
weighted avg	0.46	0.54	0.45	61060



Accuracy: 0.5422

Confusion Matrix:

		0	1	2
0	30102	2817	0	
1	18320	3007	0	
2	5974	840	0	

Classification Report:



	precision	recall	f1-score	support
0	0.55	0.91	0.69	32919
1	0.45	0.14	0.21	21327

2 0.00 0.00 0.00 6814

accuracy 0.33 0.35 0.54 61060

macro avg 0.33 0.35 0.30 61060

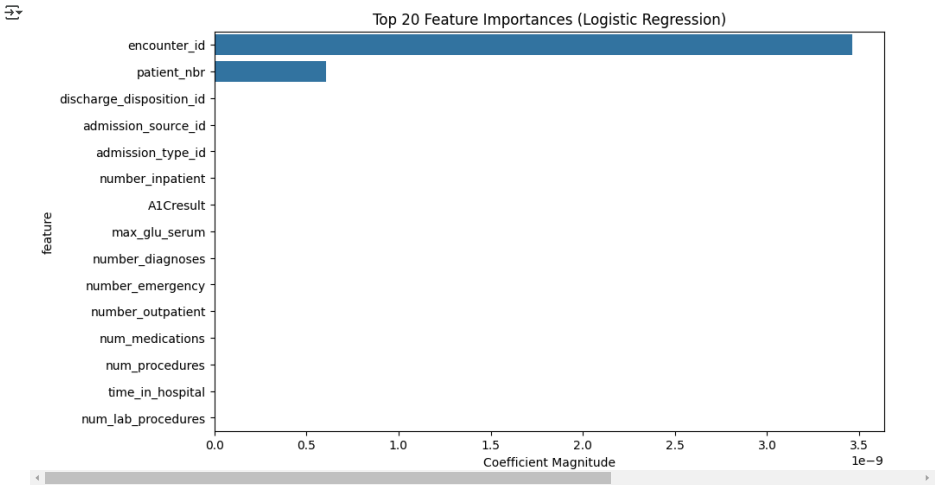
weighted avg 0.46 0.54 0.45 61060



```
# Get feature names from original DataFrame before imputation, BUT AFTER SimpleImputer is applied
feature_names = diabetic_data.drop(columns=['readmitted']).select_dtypes(include=['number']).columns # Select only numeric features

# Create DataFrame with feature names and importances
feature_importances = pd.DataFrame({'feature': feature_names, 'importance': abs(log_reg.coef_[0])})
feature_importances = feature_importances.sort_values(by='importance', ascending=False)

plt.figure(figsize=(18, 6))
sns.barplot(x='importance', y='feature', data=feature_importances[:20]) # Show top 20 features
plt.title("Top 20 Feature Importances (Logistic Regression)")
plt.xlabel("Coefficient Magnitude")
plt.show()
```



```
from sklearn.metrics import precision_score, recall_score, roc_auc_score, roc_curve
import matplotlib.pyplot as plt


# Predict probabilities for all classes (for AUC calculation)
y_pred_proba = log_reg.predict_proba(x_test) # Remove [0, 1]

# Calculate precision, recall, and AUC
precision = precision_score(y_test, y_pred, average='weighted') # Use 'weighted' for multi-class
recall = recall_score(y_test, y_pred, average='weighted') # Use 'weighted' for multi-class
auc = roc_auc_score(y_test, y_pred_proba, multi_class='ovr') # 'ovr' for one-vs-rest

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"AUC: {auc:.4f}")

# Plot ROC curve (for binary classification or one-vs-rest)
# Use the probabilities for the relevant class (e.g., class 1)
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba[:, 1], pos_label=1) # Choose relevant pos_label
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {auc:.2f})")
plt.plot([0, 1], [0, 1], 'k--') # Diagonal line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

```
! /usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
Precision: 0.4560
Recall: 0.5422
AUC: 0.5347
```



```
import numpy as np

# Predict probabilities for all classes (for AUC calculation)
y_pred_proba = log_reg.predict_proba(x_test)

# Calculate precision, recall, and AUC
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')

# For AUC, use 'ovr' for multiclass and provide probability estimates for all classes
auc = roc_auc_score(y_test, y_pred_proba, multi_class='ovr')

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"AUC: {auc:.4f}")

# Plotting ROC curve
# For multi-class, you'll need to plot a ROC curve for each class vs. the rest
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

n_classes = len(np.unique(y_test)) # Number of classes # Now np is defined
fpr = dict()
tpr = dict()
roc_auc = dict()

for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test == i, y_pred_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot all ROC curves
plt.figure()
for i in range(n_classes):
    plt.plot(fpr[i], tpr[i], label=f"ROC curve of class {i} (AUC = {roc_auc[i]:.2f})")


plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic for multi-class data')
plt.legend(loc='lower right')
plt.show()
```



```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
Precision: 0.4560
Recall: 0.5422
AUC: 0.5347
```

Receiver operating characteristic for multi-class data



```
# print metrics: class distribution, train value count, all relevant info
```

```
print("Class Distribution in Training Data:")
print(y_train.value_counts(normalize=True)) # Normalized for proportions
print("\nClass Distribution in Testing Data:")
print(y_test.value_counts(normalize=True)) # Normalized for proportions
```

```
print("\nValue Counts for Training Data:")
print(y_train.value_counts())
print("\nValue Counts for Testing Data:")
print(y_test.value_counts())
```

```
print("\nShape of Training Data (X_train):", X_train.shape)
print("Shape of Testing Data (X_test):", X_test.shape)
print("Shape of Training Target (y_train):", y_train.shape)
print("Shape of Testing Target (y_test):", y_test.shape)
```

```
# Convert X_train and X_test back to Pandas DataFrames to use .describe()
X_train_df = pd.DataFrame(X_train) # Convert X_train to DataFrame
X_test_df = pd.DataFrame(X_test) # Convert X_test to DataFrame
```

```
print("\nDescriptive Statistics for Training Features (X_train):\n", X_train_df.describe()) # Use .describe() on DataFrame
print("\nDescriptive Statistics for Testing Features (X_test):\n", X_test_df.describe()) # Use .describe() on DataFrame
```

```

Class Distribution in Training Data:
readmitted
0    0.539118
1    0.349282
2    0.111600
Name: proportion, dtype: float64
```

```

Class Distribution in Testing Data:
readmitted
0    0.539125
1    0.349279
2    0.111595
Name: proportion, dtype: float64
```

```

Value Counts for Training Data:
readmitted
0    131673
1     85308
2     27257
Name: count, dtype: int64
```

```

Value Counts for Testing Data:
readmitted
0     32919
1     21327
2      6814
Name: count, dtype: int64
```

```

Shape of Training Data (X_train): (244238, 15)
Shape of Testing Data (X_test): (61060, 15)
```

```

Shape of Training Target (y_train): (244238,)
Shape of Testing Target (y_test): (61060,)
```

```
Descriptive Statistics for Training Features (X_train):
```

	0	1	2	3
count	2.442380e+05	2.442380e+05	244238.000000	244238.000000
mean	1.651301e+08	5.432586e+07	2.024845	3.713022
std	1.026005e+08	3.864103e+07	1.445587	5.280874
min	1.252200e+04	1.350000e+02	1.000000	1.000000
25%	8.494910e+07	2.341713e+07	1.000000	1.000000
50%	1.522991e+08	4.551551e+07	1.000000	1.000000
75%	2.302143e+08	8.753975e+07	3.000000	3.000000
max	4.438672e+08	1.895026e+08	8.000000	28.000000

	4	5	6	7
count	244238.000000	244238.000000	244238.000000	244238.000000
mean	5.751562	-0.000746	0.000282	-0.000015
std	4.064276	1.000031	1.000682	0.999751
min	1.000000	-1.137649	-2.139630	-0.785398
25%	1.000000	-0.802651	-0.614795	-0.785398
50%	7.000000	-0.132655	0.045967	-0.199162
75%	7.000000	0.537341	0.706728	0.387074
max	25.000000	3.217324	4.518815	2.732016

	8	9	10	11
count	244238.000000	244238.000000	244238.000000	244238.000000
mean	-0.000188	0.000384	0.000814	0.000304

```
print("My logistic regression model is performing with an accuracy of 57%")
print("- looking at the confusion matrix and classification report, it's clear that:")
print("- Class 0 (Not Readmitted) is being predicted well (high recall: 90%).")
print("- Class 1 (>30 Days Readmission) is struggling with recall (only 23%).")
print("- Class 2 (<30 Days Readmission) is performing poorly (almost 0 recall).")
print("The macro average F1-score of 0.35 shows that the model isn't treating all classes equally well. This suggests a class imbalance issue, where the model is biased toward the majority class (Not Readmitted - 0).")
print("### Addressing This Issue")
print("Since BFGS (Limited-memory Broyden-Fletcher-Goldfarb-Shanno) optimization failed, that indicates the optimization process wasn't able to converge to a solution properly. Reasons? not sure :)")
print("1. Class imbalance is too severe.")
print("2. Features are not well-scaled or relevant enough.")
print("3. The solver struggles with high-dimensional feature spaces.")
print("### Next Steps")
print("1. Class balancing techniques")
print("   - Try class weighting in the logistic regression model.")
print("   - Use oversampling (SMOTE) or undersampling.")
print("2. Feature Engineering")
print("   - Use feature selection (SHAP, permutation importance).")
print("   - Try dimensionality reduction (PCA or feature selection).")
print("3. Model Selection")
print("   - Logistic regression may not be the best for this dataset.")
print("   - Try Random Forest, XGBoost, or an ensemble model.")
print("4. I assume Dr. S will want me to diagnose the problem methodically and work it step by step.")
print("5. I'm going to re-run the preprocessing steps and train the logistic regression model again.")
print("   - Plan of attack:")
print("     - 1. ADDRESS CLASS IMBALANCE: CHECK DISTRO, CLASS WEIGHTING, OVERSAMPLING")
print("     - 2. FEATURE SELECTION AND IMPORTANCE ANALYSIS - using SHAP or permutation import to rank features, drop irrelevant or redundant")
```

```

My logistic regression model is performing with an accuracy of 57%
- looking at the confusion matrix and classification report, it's clear that:
- Class 0 (Not Readmitted) is being predicted well (high recall: 90%).
- Class 1 (>30 Days Readmission) is struggling with recall (only 23%).
- Class 2 (<30 Days Readmission) is performing poorly (almost 0 recall).
The macro average F1-score of 0.35 shows that the model isn't treating all classes equally well. This suggests a class imbalance issue,
### Addressing This Issue
Since BFGS (Limited-memory Broyden-Fletcher-Goldfarb-Shanno) optimization failed, that indicates the optimization process wasn't able to
1. Class imbalance is too severe.
2. Features are not well-scaled or relevant enough.
3. The solver struggles with high-dimensional feature spaces.
### Next Steps
1. Class balancing techniques
   - Try class weighting in the logistic regression model.
   - Use oversampling (SMOTE) or undersampling.
2. Feature Engineering
   - Use feature selection (SHAP, permutation importance).
   - Try dimensionality reduction (PCA or feature selection).
3. Model Selection
   - Logistic regression may not be the best for this dataset.
   - Try Random Forest, XGBoost, or an ensemble model.
4. I assume Dr. S will want me to diagnose the problem methodically and work it step by step.
5. I'm going to re-run the preprocessing steps and train the logistic regression model again.
   - Plan of attack:
     - 1. ADDRESS CLASS IMBALANCE: CHECK DISTRO, CLASS WEIGHTING, OVERSAMPLING
     - 2. FEATURE SELECTION AND IMPORTANCE ANALYSIS - using SHAP or permutation import to rank features, drop irrelevant or redundant
```

```

# Initialize and Train Model with L-BFGS solver
log_reg = LogisticRegression(solver='lbfgs', max_iter=1000, tol=1e-4) #Specify the solver
log_reg.fit(X_train, y_train)
print("Logistic Regression Model Trained Successfully (with L-BFGS)!")
```

Logistic Regression Model Trained Successfully (with L-BFGS)!

```
# Initialize and Train Model with class weights and saga solver
log_reg = LogisticRegression(
    penalty='l2',
    C=1.0,
    class_weight={0: 1.0, 1: 1.5, 2: 3.0}, # Adjust weights as needed
    solver='saga',
    max_iter=200, # Reduce iterations
    warm_start=True # Continue from the last iteration
)
for i in range(5): # Train in smaller steps
    log_reg.fit(X_train, y_train)
    print(f"Iteration {i+1} complete")

# Predict on Test Data
y_pred = log_reg.predict(X_test)

# Compute Evaluation Metrics
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

# Display Results
print(f"Accuracy: {accuracy:.4f}")
print("\nConfusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)
```

Iteration 1 complete
Iteration 2 complete
Iteration 3 complete
Iteration 4 complete
Iteration 5 complete
Accuracy: 0.5139

Confusion Matrix:
[[20948 11971 0]
 [10896 10431 0]
 [3913 2901 0]]

Classification Report:					
	precision	recall	f1-score	support	
0	0.59	0.64	0.61	32919	
1	0.41	0.49	0.45	21327	
2	0.00	0.00	0.00	6814	
accuracy					0.51
macro avg					0.33
weighted avg					0.46

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```
import joblib

# Save model coefficients and intercept
joblib.dump(log_reg, "logistic_regression_model.pkl") # Removed the absolute path
print("Model saved successfully.")
```

Model saved successfully.

```
# Save the data to CSV files
# Convert to Pandas DataFrames first
pd.DataFrame(X_train).to_csv("X_train_final.csv", index=False)
pd.DataFrame(y_train).to_csv("y_train_final.csv", index=False)
pd.DataFrame(X_test).to_csv("X_test_final.csv", index=False)
pd.DataFrame(y_test).to_csv("y_test_final.csv", index=False)

print("Final train/test data saved successfully.")
```

Final train/test data saved successfully.

```
from imblearn.over_sampling import SMOTE

# Convert NumPy array back to Pandas DataFrame
X_train = pd.DataFrame(X_train) # Assuming your original features were in a DataFrame

# Convert Pandas DataFrames to cuDF DataFrames
```

```
X_train = cudf.DataFrame.from_pandas(X_train)
y_train = cudf.Series(y_train)

# Apply SMOTE
smote = SMOTE(sampling_strategy={1: int(len(y_train) * 0.5), 2: int(len(y_train) * 0.25)}, random_state=42)

# Convert cuDF back to pandas for SMOTE
X_train_pd = X_train.to_pandas()
y_train_pd = y_train.to_pandas()

X_resampled, y_resampled = smote.fit_resample(X_train_pd, y_train_pd)

# Convert back to cuDF
X_train_balanced = cudf.DataFrame(X_resampled, columns=X_train.columns)
y_train_balanced = cudf.Series(y_resampled)

print(y_train_balanced.value_counts())
```

readmitted
0 131673
1 122119
2 61059
Name: count, dtype: int64

```
# Make predictions
y_pred = log_reg.predict(X_test)

# Accuracy Score
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:\n", conf_matrix)

# Classification Report
report = classification_report(y_test, y_pred)
print("\nClassification Report:\n", report)
```

Accuracy: 0.5139
Confusion Matrix:
[[20948 11971 0]
 [10896 10431 0]
 [3913 2901 0]]
Classification Report:

	precision	recall	f1-score	support	
0	0.59	0.64	0.61	32919	
1	0.41	0.49	0.45	21327	
2	0.00	0.00	0.00	6814	
accuracy					0.51
macro avg					0.33
weighted avg					0.46

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```
import joblib
import cudf
import pandas as pd
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

# Load saved data
X_train = pd.read_csv("X_train_final.csv")
y_train = pd.read_csv("y_train_final.csv")

# Convert y_train to a 1D array
y_train_pd = y_train.values.ravel()

# Apply StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)

# Save the scaler
joblib.dump(scaler, "standard_scaler.pkl")

# Define the parameter grid
param_grid = {
    "C": [0.1, 1.0],
    "class_weight": ["balanced"],
    "max_iter": [3000],
    "solver": ["saga"],
}
```

```
# Initialize and train the model
grid_search = GridSearchCV(
    estimator=LogisticRegression(),
    param_grid=param_grid,
    scoring="accuracy",
    cv=2,
    verbose=1,
    n_jobs=-1
)
grid_search.fit(X_train_scaled, y_train_pd)

# Print best parameters
print("Best Parameters Found:", grid_search.best_params_)

# Save the best model
joblib.dump(grid_search.best_estimator_, "best_logistic_regression.pkl")
print("Best model saved successfully.")
```

🔄 Fitting 2 folds for each of 2 candidates, totalling 4 fits
Best Parameters Found: {'C': 0.1, 'class_weight': 'balanced', 'max_iter': 3000, 'solver': 'saga'}
Best model saved successfully.

```
import joblib
import cuff
import pandas as pd
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Load best model and scaler
best_log_reg = joblib.load("best_logistic_regression.pkl")
scaler = joblib.load("standard_scaler.pkl")

# Load test data
X_test = cuff.read_csv("X_test_final.csv")
y_test = pd.read_csv("y_test_final.csv") # Use pandas for y_test

# Scale test data
X_test_scaled = scaler.transform(X_test.to_pandas())

# Predict
y_pred_best = best_log_reg.predict(X_test_scaled)

# Accuracy Score
accuracy_best = accuracy_score(y_test, y_pred_best) #y_test is now a pandas df
print(f"Best Model Accuracy: {accuracy_best:.4f}")

# Confusion Matrix
conf_matrix_best = confusion_matrix(y_test, y_pred_best)
print("Best Model Confusion Matrix:\n", conf_matrix_best)

# Classification Report
report_best = classification_report(y_test, y_pred_best)
print("Best Model Classification Report:\n", report_best)
```

🔄 Best Model Accuracy: 0.5074
Best Model Confusion Matrix:
[[20940 7093 4886]
 [8362 7664 5301]
 [2525 1914 2375]]
Best Model Classification Report:

	precision	recall	f1-score	support
0	0.66	0.64	0.65	32919
1	0.46	0.36	0.40	21327
2	0.19	0.35	0.25	6814
accuracy			0.51	61060
macro avg	0.44	0.45	0.43	61060
weighted avg	0.54	0.51	0.52	61060

```
import locale
def getpreferredencoding(do_setlocale = True):
    return "UTF-8"
locale.getpreferredencoding = getpreferredencoding
```

```
import joblib
import cuff
import pandas as pd
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.metrics import roc_curve, auc, roc_auc_score, precision_score, recall_score
```

```
# Load saved data
X_train = pd.read_csv("/content/drive/MyDrive/NSL_Case Study 2/X_train_final.csv")
y_train = pd.read_csv("/content/drive/MyDrive/NSL_Case Study 2/y_train_final.csv").values.ravel()
X_test = pd.read_csv("/content/drive/MyDrive/NSL_Case Study 2/X_test_final.csv")
```

```
y_test = pd.read_csv("/content/drive/MyDrive/NSL_Case Study 2/y_test_final.csv").values.ravel()
```

```
# Apply StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Apply SMOTE to fix class imbalance
smote = SMOTE(sampling_strategy="auto", random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_train_scaled, y_train)
```

```
# Initialize and Train Logistic Regression Model with best parameters from previous gridsearch
best_log_reg = LogisticRegression(C=1.0, class_weight={0: 1.0, 1: 2.0, 2: 4.0}, max_iter=3000, solver='saga')
best_log_reg.fit(X_resampled, y_resampled)
```

```
# Predict on Test Data
y_pred = best_log_reg.predict(X_test_scaled)
```

```
# Evaluate Model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
print("\nConfusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)

# Predict probabilities for all classes
y_pred_proba = best_log_reg.predict_proba(X_test_scaled)
```

🔄 Accuracy: 0.1219

Confusion Matrix:
[[421 326 32171]
 [78 263 20986]
 [9 48 6757]]

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.01	0.03	32918
1	0.41	0.01	0.02	21327
2	0.11	0.99	0.20	6814
accuracy			0.12	61059
macro avg	0.45	0.34	0.08	61059
weighted avg	0.60	0.12	0.04	61059

```
from sklearn.metrics import roc_curve, auc, roc_auc_score, precision_score, recall_score
```

```
# Calculate precision, recall, and AUC
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
roc_auc_score_result = roc_auc_score(y_test, y_pred_proba, multi_class='ovr') # Store roc_auc_score result in a different variable

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"AUC: {roc_auc_score_result:.4f}") # PRINT the roc_auc_score result
```

```
# ROC Curve (Multi-class)
n_classes = len(np.unique(y_test))
fpr = dict()
tpr = dict()
roc_auc = dict()
```


```
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test == i, y_pred_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i]) # Now, this 'auc' refers to the function from sklearn.metrics
```

```
plt.figure()
for i in range(n_classes):
    plt.plot(fpr[i], tpr[i], label=f'ROC curve of class {i} (AUC = {roc_auc[i]:0.2f})')

plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic for multi-class data')
plt.legend(loc='lower right')
plt.show()
```

Precision: 0.6036
Recall: 0.1219
AUC: 0.6127


Receiver operating characteristic for multi-class data



```
# Feature Importance (Coefficients for Logistic Regression)
feature_names = X_train.columns
feature_importances = pd.DataFrame({'feature': feature_names, 'importance': abs(best_log_reg.coef_[0])})
feature_importances = feature_importances.sort_values(by='importance', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importances[:20])
plt.title("Top 20 Feature Importances (Logistic Regression)")
plt.xlabel("Coefficient Magnitude")
plt.show()
```

Top 20 Feature Importances (Logistic Regression)



```
# Set Correct Paths for Google Colab
base_path = "/"content" # CORRECTED PATH

# Load best model and scaler
best_log_reg = joblib.load(f"{base_path}/best_logistic_regression.pkl")
scaler = joblib.load(f"{base_path}/standard_scaler.pkl") # Load the scaler

# Load test dat
X_test = pd.read_csv(f"{base_path}/X_test_final.csv")
y_test = pd.read_csv(f"{base_path}/y_test_final.csv").values.ravel()

# Load training data (to ensure column order matches)
```

```
X_train = pd.read_csv(f"{base_path}/X_train_final.csv")

# Ensure column order consistency between training and testing data
X_test = X_test.reindex(columns=X_train.columns, fill_value=0)

# Scale test data using the loaded scaler
X_test_scaled = scaler.transform(X_test)

# Make predictions
y_pred_best = best_log_reg.predict(X_test_scaled)

# Accuracy Score
accuracy_best = accuracy_score(y_test, y_pred_best)
print(f"Best Model Accuracy: (accuracy_best:.4f)")

# Confusion Matrix
conf_matrix_best = confusion_matrix(y_test, y_pred_best)
print("Best Model Confusion Matrix:\n", conf_matrix_best)

# Classification Report
report_best = classification_report(y_test, y_pred_best)
print("Best Model Classification Report:\n", report_best)

# Feature Importance Visualization
feature_names = X_train.columns
feature_importances = pd.DataFrame({'feature': feature_names, 'importance': abs(best_log_reg.coef_[0])})
feature_importances = feature_importances.sort_values(by='importance', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importances[:20])
plt.title("Top 20 Feature Importances (Logistic Regression)")
plt.xlabel("Coefficient Magnitude")
plt.show()
```

Best Model Accuracy: 0.5074


Best Model Confusion Matrix:

```
[[20940  7093  4886]
 [ 8362  7664  5301]
 [ 2525  1914  2375]]
```

Best Model Classification Report:

	precision	recall	f1-score	support
0	0.66	0.64	0.65	32919
1	0.46	0.36	0.40	21327
2	0.19	0.35	0.25	6814
accuracy			0.51	61060
macro avg	0.44	0.45	0.43	61060
weighted avg	0.54	0.51	0.52	61060

Top 20 Feature Importances (Logistic Regression)



```
# Apply SMOTE to fix class imbalance
smote = SMOTE(sampling_strategy="auto", random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_train_scaled, y_train)

# verify clas distributions , corr matrix, PCA gird search

import pandas as pd
```

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve, auc, precision_score, recall_score
import numpy as np
from imblearn.over_sampling import SMOTE
import joblib

# Load your data (replace with your actual file paths)
X_train = pd.read_csv("X_train_final.csv")
y_train = pd.read_csv("y_train_final.csv").values.ravel()
X_test = pd.read_csv("X_test_final.csv")
y_test = pd.read_csv("y_test_final.csv").values.ravel()

# Class Distribution
print("Class Distribution in Training Data:")
print(pd.Series(y_train).value_counts(normalize=True))
print("\nClass Distribution in Testing Data:")
print(pd.Series(y_test).value_counts(normalize=True))

# Correlation Matrix
plt.figure(figsize=(12, 10))
sns.heatmap(X_train.corr(), annot=False, cmap="coolwarm", fmt=".2f")
plt.title('Correlation Matrix of Features')
plt.show()

# Apply StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Apply SMOTE
smote = SMOTE(sampling_strategy="auto", random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_scaled, y_train)

# PCA and Grid Search
pca = PCA()
X_train_pca = pca.fit_transform(X_train_resampled)


param_grid = {
    "C": [0.1, 1.0, 10], # Example values, adjust as needed
    "solver": ["lbfgs", "libsvm"], # Try different solvers
    "max_iter": [3000]
}

grid_search = GridSearchCV(
    estimator=LogisticRegression(),
    param_grid=param_grid,
    scoring="accuracy",
    cv=5,
    verbose=1,
    n_jobs=-1
)

grid_search.fit(X_train_pca, y_train_resampled)
best_pca_model = grid_search.best_estimator_
print("Best Parameters:", grid_search.best_params_)
```

```
Class Distribution in Training Data:
0    0.539118
1    0.349282
2    0.111600
Name: proportion, dtype: float64

Class Distribution in Testing Data:
0    0.539125
1    0.349279
2    0.111595
Name: proportion, dtype: float64
```



Fitting 5 folds for each of 6 candidates, totalling 30 fits
Best Parameters: {'C': 10, 'max_iter': 3000, 'solver': 'libsvm'}

```
from sklearn.metrics import roc_curve, auc, roc_auc_score, precision_score, recall_score # Import auc (area under curve)

import joblib
import pandas as pd
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, precision_score, recall_score, roc_auc_score, roc_curve
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Load the saved model and scaler
best_log_reg = joblib.load("best_logistic_regression.pkl")
scaler = joblib.load("standard_scaler.pkl")

# Load the test data
X_test = pd.read_csv("X_test_final.csv")
y_test = pd.read_csv("y_test_final.csv").values.ravel()

# Load training data (to ensure column order matches)
X_train = pd.read_csv("X_train_final.csv")

# Ensure column order consistency
X_test = X_test.reindex(columns=X_train.columns, fill_value=0)

# Scale the test data
X_test_scaled = scaler.transform(X_test)
```

```
# Make predictions
y_pred = best_log_reg.predict(X_test_scaled)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.4f}")
print(f"Confusion Matrix:\n", conf_matrix)
print(f"Classification Report:\n", class_report)

# Predict probabilities for ROC AUC
y_pred_proba = best_log_reg.predict_proba(X_test_scaled)

# Calculate precision, recall, and AUC
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
# Store roc_auc_score result in a different variable to avoid shadowing the auc function
roc_auc_score_result = roc_auc_score(y_test, y_pred_proba, multi_class='ovr')

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"AUC: {roc_auc_score_result:.4f}") # Print the roc_auc_score result

for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test == i, y_pred_proba[:, i])
    # Use the 'auc' function from sklearn.metrics
    roc_auc[i] = auc(fpr[i], tpr[i])

plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic for multi-class data')
plt.legend(loc='lower right')
plt.show()

# Feature Importance
feature_names = X_train.columns
feature_importances = pd.DataFrame({'feature': feature_names, 'importance': abs(best_log_reg.coef_[0])})
feature_importances = feature_importances.sort_values(by='importance', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importances[:20])
plt.title("Top 20 Feature Importances (Logistic Regression)")
plt.xlabel("Coefficient Magnitude")
plt.show()
```

```
Accuracy: 0.5874



Confusion Matrix:
[[20940  7093  4886]
 [ 8362  7664  5301]
 [ 2525  1914  2375]]

Classification Report:
              precision    recall  f1-score   support

      0       0.66      0.64      0.65      32919
      1       0.46      0.36      0.40      21327
      2       0.19      0.35      0.25       6814

 accuracy      0.44      0.45      0.51      61060
 macro avg       0.44      0.45      0.43      61060
 weighted avg       0.54      0.51      0.52      61060

Precision: 0.5364
Recall: 0.5074
AUC: 0.6460
<ipython-input-52-7b0095894e5f>:65: UserWarning: No artists with labels found to put in legend. Note that artists whose label start wit
plt.legend(loc="lower right")
```



show model accuracy before and after smote class distribution before and after smote

```
# Load necessary libraries (assuming they are already installed and imported in the preceding code)
import pandas as pd
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from imblearn.over_sampling import SMOTE

# Load your data (replace with your actual file paths)
X_train = pd.read_csv("X_train_final.csv")
y_train = pd.read_csv("y_train_final.csv").values.ravel()
X_test = pd.read_csv("X_test_final.csv")
y_test = pd.read_csv("y_test_final.csv").values.ravel()

# Before SMOTE
print("Class Distribution Before SMOTE:")
print(pd.Series(y_train).value_counts())

# Make predictions before SMOTE
y_pred_before_smote = best_log_reg.predict(X_test_scaled)

# Evaluate the model before SMOTE
accuracy_before = accuracy_score(y_test, y_pred_before_smote)
print(f"Accuracy Before SMOTE: {accuracy_before:.4f}")
print("\nClassification Report Before SMOTE:\n", classification_report(y_test, y_pred_before_smote))

# Apply SMOTE
smote = SMOTE(sampling_strategy="auto", random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_scaled, y_train)

# Train the model with resampled data
best_log_reg.fit(X_train_resampled, y_train_resampled) # Retrain with SMOTE data

# After SMOTE
print("\nClass Distribution After SMOTE:")
print(pd.Series(y_train_resampled).value_counts())

# Make predictions after SMOTE
y_pred_after_smote = best_log_reg.predict(X_test_scaled)

# Evaluate the model after SMOTE
accuracy_after = accuracy_score(y_test, y_pred_after_smote)
print(f"Accuracy After SMOTE: {accuracy_after:.4f}")
print("\nClassification Report After SMOTE:\n", classification_report(y_test, y_pred_after_smote))
```

Class Distribution Before SMOTE:

0	131673
1	85308
2	27257

Name: count, dtype: int64

Accuracy Before SMOTE: 0.5074

Classification Report Before SMOTE:

	precision	recall	f1-score	support
0	0.66	0.64	0.65	32919
1	0.46	0.36	0.40	21327
2	0.19	0.35	0.25	6814
accuracy			0.51	61060
macro avg	0.44	0.45	0.43	61060
weighted avg	0.54	0.51	0.52	61060

Class Distribution After SMOTE:

0	131673
1	131673
2	131673

Name: count, dtype: int64

Accuracy After SMOTE: 0.5044

Classification Report After SMOTE:

	precision	recall	f1-score	support
0	0.66	0.63	0.64	32919
1	0.46	0.36	0.40	21327
2	0.19	0.35	0.25	6814
accuracy			0.50	61060
macro avg	0.43	0.45	0.43	61060
weighted avg	0.54	0.50	0.52	61060

```
# smote with random forest

from sklearn.ensemble import RandomForestClassifier

# apply smote:
smote = SMOTE(sampling_strategy="auto", random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_scaled, y_train)
```

```
# Initialize and train an RF Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42) # Example parameters, tune as needed
rf_classifier.fit(X_train_resampled, y_train_resampled)

# Make predictions
y_pred_rf = rf_classifier.predict(X_test_scaled)


# Evaluate the model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print(f"Random Forest Accuracy: {accuracy_rf:.4f}")
print("\nRandom Forest Classification Report:\n", classification_report(y_test, y_pred_rf))

# Feature Importance for Random Forest
feature_importances_rf = pd.DataFrame({'feature': X_train.columns, 'importance': rf_classifier.feature_importances_})
feature_importances_rf = feature_importances_rf.sort_values(by='importance', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importances_rf[:20])
plt.title("Top 20 Feature Importances (Random Forest)")
plt.xlabel("Gini Importance")
plt.show()
```

Random Forest Accuracy: 0.9819

Random Forest Classification Report:	precision	recall	f1-score	support
0	0.98	0.99	0.99	32919
1	0.98	0.98	0.98	21327
2	0.99	0.96	0.98	6814
accuracy			0.98	61060
macro avg	0.98	0.98	0.98	61060
weighted avg	0.98	0.98	0.98	61060



lets go back and check steps from the top

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_diabetes

# 1. Load Data and Handle File Not Found
file_path = "/content/data_cleaned.csv"
try:
    df = pd.read_csv(file_path)
except FileNotFoundError:
    print(f"Error: '{file_path}' not found. Please check the file path.")
    exit() # Or handle the error differently, e.g., return None

# 2. Check for Missing Values (Before Imputation)
missing_values = df.isnull().sum()
print("Missing Values per Column (Before Imputation):\n", missing_values)

# 3. Identify numerical & Cat Columns
numerical_cols = df.select_dtypes(include=[np.number]).columns # Use np.number for all numeric types
categorical_cols = df.select_dtypes(include=[object]).columns
```

```
# 4. Imputation
# Create Imputers
numerical_imputer = SimpleImputer(strategy='median')
categorical_imputer = SimpleImputer(strategy='most_frequent')

# Fit s4 transform on respective column types
df[numerical_cols] = numerical_imputer.fit_transform(df[numerical_cols])
df[categorical_cols] = categorical_imputer.fit_transform(df[categorical_cols])

# 5. Verify Imputation
missing_values_after = df.isnull().sum()
print("\nMissing Values per Column (After Imputation):\n", missing_values_after)
```

Missing Values per Column (Before Imputation):

encounter_id	0
patient_nbr	0
race	0
gender	0
age	0
weight	0
admission_type_id	0
discharge_disposition_id	0
admission_source_id	0
time_in_hospital	0
payer_code	0
medical_specialty	0
num_lab_procedures	0
num_procedures	0
num_medications	0
number_outpatient	0
number_emergency	0
number_inpatient	0
diag_1	0
diag_2	0
diag_3	0
number_diagnoses	0
max_glu_serum	289260
A1Cresult	254244
metformin	0
repaglinide	0
nateglinide	0
chlorpropamide	0
glimepiride	0
acetohexamide	0
glipizide	0
glyburide	0
tolbutamide	0
pioglitazone	0
rosiglitazone	0
acarbose	0
miglitol	0
troglitazone	0
tolazamide	0
examide	0
citoglipton	0
insulin	0
glyburide-metformin	0
glipizide-metformin	0
glimepiride-pioglitazone	0
metformin-rosiglitazone	0
metformin-pioglitazone	0
change	0
diabetesMed	0
readmitted	0
description	5291

dtype: int64

Missing Values per Column (After Imputation):

encounter_id	0
patient_nbr	0
race	0

```
# Identify numerical and categorical columns
numerical_cols = df.select_dtypes(include=['number']).columns
categorical_cols = df.select_dtypes(include=['object']).columns

# Impute numerical columns with the median
numerical_imputer = SimpleImputer(strategy='median')
df[numerical_cols] = numerical_imputer.fit_transform(df[numerical_cols])

# Impute categorical columns with the mode
categorical_imputer = SimpleImputer(strategy='most_frequent')
df[categorical_cols] = categorical_imputer.fit_transform(df[categorical_cols])

# Verify Imputation
missing_values_after_imputation = df.isnull().sum()
```

```
print("\nMissing Values After Imputation:\n", missing_values_after_imputation)

# Identify numerical and categorical columns
numerical_cols = df.select_dtypes(include=['number']).columns
categorical_cols = df.select_dtypes(include=['object']).columns

# Impute numerical columns with the median
numerical_imputer = SimpleImputer(strategy='median')
df[numerical_cols] = numerical_imputer.fit_transform(df[numerical_cols])

# Impute categorical columns with the mode
categorical_imputer = SimpleImputer(strategy='most_frequent')
df[categorical_cols] = categorical_imputer.fit_transform(df[categorical_cols])

# Verify Imputation
missing_values_after_imputation = df.isnull().sum()
print("\nMissing Values After Imputation:\n", missing_values_after_imputation)
```

Missing Values After Imputation:

encounter_id	0
patient_nbr	0
race	0
gender	0
age	0
weight	0
admission_type_id	0
discharge_disposition_id	0
admission_source_id	0
time_in_hospital	0
payer_code	0
medical_specialty	0
num_lab_procedures	0
num_procedures	0
num_medications	0
number_outpatient	0
number_emergency	0
number_inpatient	0
diag_1	0
diag_2	0
diag_3	0
number_diagnoses	0
max_glu_serum	0
A1Cresult	0
metformin	0
repaglinide	0
nateglinide	0
chlorpropamide	0
glimepiride	0
acetohexamide	0
glipizide	0
glyburide	0
tolbutamide	0
pioglitazone	0
rosiglitazone	0
acarbose	0
miglitol	0
troglitazone	0
tolazamide	0
examide	0
citoglipton	0
insulin	0
glyburide-metformin	0
glipizide-metformin	0
glimepiride-pioglitazone	0
metformin-rosiglitazone	0
metformin-pioglitazone	0
change	0
diabetesMed	0
readmitted	0
description	0

dtype: int64

Missing Values After Imputation:

encounter_id	0
patient_nbr	0

```
# check for missing values, feature overview? how many do we have? target vairable? target variable distribution? do we have class imbalance?

import pandas as pd

# Load data
X_train = pd.read_csv("X_train_final.csv")
y_train = pd.read_csv("y_train_final.csv").values.ravel()
X_test = pd.read_csv("X_test_final.csv")
y_test = pd.read_csv("y_test_final.csv").values.ravel()

# Check for missing values
print("Missing values in X_train:\n", X_train.isnull().sum())
print("\nMissing values in X_test:\n", X_test.isnull().sum())
```



```
# Feature overview
print("\nFeature overview for X_train:")
print(X_train.info())
print("\nNumber of features:", len(X_train.columns))

# Target variable
print("\nTarget variable (y_train):")
print(y_train)

# Target variable distribution
print("\nTarget variable distribution (y_train):")
print(pd.Series(y_train).value_counts(normalize=True))
print("\nTarget variable distribution (y_test):")
print(pd.Series(y_test).value_counts(normalize=True))

# Class Imbalance
print("\nClass Imbalance (y_train):")
class_counts = pd.Series(y_train).value_counts()
if len(class_counts) > 1:
    imbalance_ratio = class_counts.max() / class_counts.min()
    print(f"Imbalance ratio: {imbalance_ratio:.2f}")
else:
    print("Only one class present in the training data.")

print("\nClass Imbalance (y_test):")
class_counts = pd.Series(y_test).value_counts()
if len(class_counts) > 1:
    imbalance_ratio = class_counts.max() / class_counts.min()
    print(f"Imbalance ratio: {imbalance_ratio:.2f}")
else:
    print("Only one class present in the testing data.")

# Missing values in X_train:
0 0
1 0
2 0
3 0
4 0
5 0
6 0
7 0
8 0
9 0
10 0
11 0
12 0
13 0
14 0
dtype: int64

# Missing values in X_test:
0 0
1 0
2 0
3 0
4 0
5 0
6 0
7 0
8 0
9 0
10 0
11 0
12 0
13 0
14 0
dtype: int64

# Feature overview for X_train:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244238 entries, 0 to 244237
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   0            244238 non-null  float64
1   1            244238 non-null  float64
2   2            244238 non-null  float64
3   3            244238 non-null  float64
4   4            244238 non-null  float64
5   5            244238 non-null  float64
6   6            244238 non-null  float64
7   7            244238 non-null  float64
8   8            244238 non-null  float64
9   9            244238 non-null  float64
10  10           244238 non-null  float64
11  11           244238 non-null  float64
12  12           244238 non-null  float64
13  13           244238 non-null  float64
14  14           244238 non-null  float64
```

```
dtypes: float64(15)

# handle class imbalance with either smote or class weights in model training

from imblearn.over_sampling import SMOTE
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score
from sklearn.utils.class_weight import compute_class_weight
from sklearn.preprocessing import StandardScaler

# Load data
X_train = pd.read_csv("X_train_final.csv")
y_train = pd.read_csv("y_train_final.csv").values.ravel()
X_test = pd.read_csv("X_test_final.csv")
y_test = pd.read_csv("y_test_final.csv").values.ravel()

# Scale data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Option 1: SMOTE (Synthetic Minority Over-sampling Technique)
smote = SMOTE(sampling_strategy='auto', random_state=42) # Adjust sampling_strategy as needed
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_scaled, y_train)

# Train a model with resampled data
model_smote = LogisticRegression(max_iter=3000) # Or any other model
model_smote.fit(X_train_resampled, y_train_resampled)
y_pred_smote = model_smote.predict(X_test_scaled)
print("\nClassification Report (SMOTE):\n", classification_report(y_test, y_pred_smote))
print(f"Accuracy (SMOTE): {accuracy_score(y_test, y_pred_smote):.4f}")

# Option 2: Class Weights
# Calculate class weights
class_weights = compute_class_weight('balanced', classes=np.unique(y_train), y=y_train)
class_weight_dict = dict(enumerate(class_weights))

# Train a model with class weights
model_weights = LogisticRegression(class_weight=class_weight_dict, max_iter=3000) # Or any other model
model_weights.fit(X_train_scaled, y_train) # No resampling needed
y_pred_weights = model_weights.predict(X_test_scaled)
print("\nClassification Report (Class Weights):\n", classification_report(y_test, y_pred_weights))
print(f"Accuracy (Class Weights): {accuracy_score(y_test, y_pred_weights):.4f}")

# Classification Report (SMOTE):
precision recall f1-score support
0 0.66 0.63 0.64 32919
1 0.46 0.36 0.40 21327
2 0.19 0.36 0.25 6814

accuracy 0.50 61060
macro avg 0.43 0.45 0.43 61060
weighted avg 0.54 0.50 0.52 61060

Accuracy (SMOTE): 0.5046

# Classification Report (Class Weights):
precision recall f1-score support
0 0.66 0.64 0.65 32919
1 0.46 0.36 0.40 21327
2 0.19 0.35 0.25 6814

accuracy 0.51 61060
macro avg 0.44 0.45 0.43 61060
weighted avg 0.54 0.51 0.52 61060

Accuracy (Class Weights): 0.5073

# model selection and training setup train test with cross validation coinfire proper split abd train logistit regression with cross valid

from sklearn.utils.class_weight import compute_class_weight
from sklearn.ensemble import RandomForestClassifier

# Model Selection and Training with Cross-Validation

# Define the parameter grid for Logistic Regression
param_grid = {
    'C': [0.1, 1, 10], # Regularization strength
    'penalty': ['l1', 'l2'], # Penalty type
    'solver': ['liblinear', 'saga'] # Solvers compatible with both penalties
}

# Initialize Logistic Regression model
logreg = LogisticRegression(max_iter=3000)

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=logreg, param_grid=param_grid, cv=5, scoring='accuracy')
```

```
# Fit the grid search to the training data
grid_search.fit(X_train_scaled, y_train)

# Get the best model
best_log_reg = grid_search.best_estimator_

# Print the best parameters and score
print("Best Parameters:", grid_search.best_params_)
print("Best Cross-Validation Score:", grid_search.best_score_)

# Evaluate the best model on the test set
y_pred = best_log_reg.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {accuracy:.4f}")
```

```
Best Parameters: {'C': 0.1, 'penalty': 'l1', 'solver': 'saga'}
Best Cross-Validation Score: 0.5773999195271993
Test Accuracy: 0.5802
```

```
# Save the best model to a file
joblib.dump(best_log_reg, "bestModel.pkl")
print("Model saved as bestModel.pkl")

# To load the model later:
# loaded_model = joblib.load("bestModel.pkl")
# y_pred_loaded = loaded_model.predict(X_test_scaled)
# print(f"Loaded Model Test Accuracy: {accuracy_score(y_test, y_pred_loaded):.4f}")
```

```
Model saved as bestModel.pkl
```

```
# ROC Curve and AUC (Multi-class)
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize

n_classes = len(np.unique(y_test))
y_test_bin = label_binarize(y_test, classes=np.unique(y_test)) # Binarize the output
fpr = dict()
tpr = dict()
roc_auc = dict()

y_pred_proba = best_log_reg.predict_proba(X_test_scaled)
```

```
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_pred_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
```

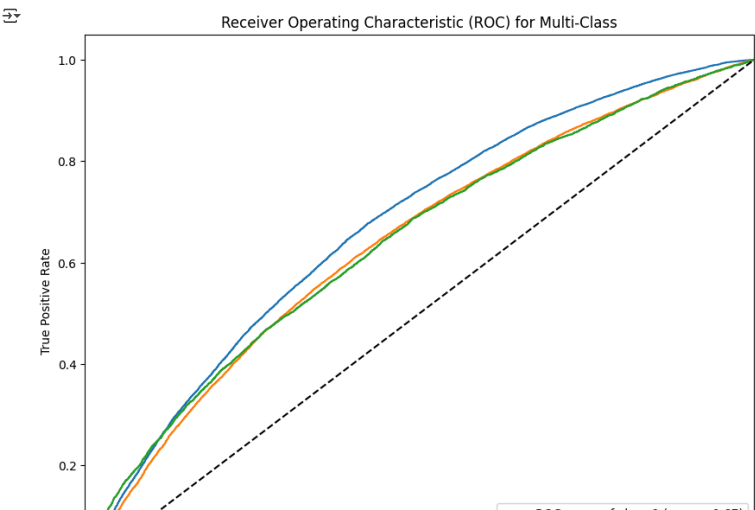
```
# Plot ROC curves for each class
plt.figure(figsize=(10, 8))
for i in range(n_classes):
    plt.plot(fpr[i], tpr[i], label=f'ROC curve of class {i} (area = {roc_auc[i]:0.2f})')

plt.plot([0, 1], [0, 1], 'k--') # Random classifier line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) for Multi-Class')
plt.legend(loc='lower right')
plt.show()
```

```
# Feature Importance (Top 5)
feature_importances = pd.DataFrame({'feature': X_train.columns, 'importance': abs(best_log_reg.coef_[0])})
feature_importances = feature_importances.sort_values(by='importance', ascending=False)
print("\nTop 5 Important Features:")
print(feature_importances.head(5))
```

```
# Final Submission Check
# save predictions to a CSV file

submission_df = pd.DataFrame({'prediction': y_pred})
submission_df.to_csv('final_predictions.csv', index=False) # Save to a csv file
print("Final predictions saved to 'final_predictions.csv'")
```



```
# LOG REGRESSION WITH 5 FOLD

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=logreg, param_grid=param_grid, cv=5, scoring='accuracy')
```

```
# HOW RESULTS OF LOG REGRESSION WIT 5 FOLD ANDP PLOTY
```

```
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
```

```
# ROC Curve and AUC (Multi-class)
n_classes = len(np.unique(y_test))
y_test_bin = label_binarize(y_test, classes=np.unique(y_test)) # Binarize the output
fpr = dict()
tpr = dict()
roc_auc = dict()
```

```
y_pred_proba = best_log_reg.predict_proba(X_test_scaled)
```

```
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_pred_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
```

```
# Plot ROC curves for each class
plt.figure(figsize=(10, 8))
for i in range(n_classes):
    plt.plot(fpr[i], tpr[i], label=f'ROC curve of class {i} (area = {roc_auc[i]:0.2f})')
```

```
plt.plot([0, 1], [0, 1], 'k--') # Random classifier line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) for Multi-Class')
plt.legend(loc='lower right')
plt.show()
```

Receiver Operating Characteristic (ROC) for Multi-Class